CSCI-UA.0480-003
Parallel Computing
Lab 3

You will be provided with a sequential program maxseq.c that can be compiled and run with:

gcc -o maxseq maxseq.c

Then you execute it with:

./maxseq x

where x is a positive number.

The program then will allocate an array of x elements of unsigned int, fill it with random numbers from 0 to x and then call a function getmax that returns the maximum number of the array. Several members of the array may have that largest number.

Coding:

A. Keep a copy of the sequential version maxseq.c
B. Make another copy of the above file and rename it maxgpu.cu
C. Modify that file maxseq.cu such that instead of calling getmax you call a kernel getmaxcu() that outsources the job of finding the maximum number to the GPU. Let the main function allocate and populate the array in the host. Then you need to allocate memory in the device for that array and transfer the array from the host to the device, calculate the maximum in the device, then transfer that number back to the host.
D. Let the printf() in the host print the maximum number as in the sequential version.
E. Measure the time for both the sequential and the CUDA versions.

Reporting Results (the report):

Put one line at the beginning stating which machine have you used to do the experiments (cuda1, cuda2, …).

1. Explain how you picked the block and grid sizes/dimentions, and justify your choices.
2. Write the command line you used to compile the CUDA version.
3. Draw a bar graph as follows:
    a. The Y-axis is the time
    b. X-axis is the length of the array (i.e. the number of elements), try with 1000, 10000, 100000, 1000000 (one million), 10000000 (ten million), and 100000000 (hundred million)
    c. For each value of X show two bars: the sequential version and the CUDA version
4. Explain the behavior in the graph.

What to submit through NYU classes
(in single zip file called yourlasname.firstname.zip):
- the source code maxgpu.cu
- the report in pdf

Grading policy:
- maxgpu.cu compiles and run correctly 10 points. Otherwise, we will look at the code and give you a partial credit that does not exceed 20/50
- #1 in the reporting results above: 10 points (Let's see how you design your project)
- #3 15 points
- #4 15 points
- If your code does not compile and yet you show results, or if you give us a code that does not have CUDA code in it, you will get a zero.

Have Fun!